CHRISTINE NGUYEN

Head of AI Enablement | Principal, AI Transformation, Value & Portfolio | Intake → Prioritization → Adoption, with Governed Agentic Capability

Austin, TX linkedin.com/in/christinenguyen7 christineqnguyen@gmail.com christinenguyen.me · github.com/christine-q-nguyen

SUMMARY

Operating leader for enterprise AI enablement: I run intake, prioritization, and portfolio governance as one system, author the operating model, and turn pilots into leadership decisions and adopted capability. Fourteen months running firm-wide GenAI enablement at McKinsey (344 opportunities, 23 practices) on top of 18 years leading product, design, and organizational change. **My perspective on GenAI and technology is simple: it only matters if it fits the work, the people, and the operating model.**

CORE CAPABILITIES

AI TRANSLATION	Ambiguous business needs into AI opportunities, requirements, and adoption plans; intake as a system; DFV/RICE triage; decision logs; build-vs-buy; demand taxonomies	**LEARNING DESIGN**	Workshops, office hours, hackathons, role-based coaching, co-build then hand off, design-thinking instruction, AI adoption support
LEADERSHIP	Executive storytelling, operating-model design, decision-ready narratives, leadership forums, global team coaching, distributed work planning	**WORKFLOW DESIGN**	Recorded workflow capture, stakeholder interviews, think-aloud testing, service blueprinting, journey mapping, process redesign, JTBD PRDs
RESPONSIBLE AI	Human review gates, confidence thresholds, binary pass/fail rubrics, LLM-judge validation, risk–coverage curves, source attribution, least privilege, data-use constraints	**AI PROTOTYPING**	Agentic workflows and skills, prompt architecture, structured outputs, transcript parsing, MCP clients and servers, Python (CLI, eval harness), Jira orchestration

EXPERIENCE

Principal Product Manager – GenAI Transformation & Enablement Lead – Apr 2025 – Present

McKinsey & Company

Platform AI Advisory & Integration | previously Product Manager, Practice Technology AI Programs | Austin, TX

- Run firm-wide GenAI intake and prioritization as a system across 23 practices, internal functions, and a research institute: a 344-opportunity board (Oct 2026) triaged with DFV scoring and a start/stop/continue/defer decision log; in my first year I opened 68 of 134 new opportunities (51%).
- Run the forums that turn pilots into decisions (weekly practice syncs, bi-weekly leadership connects, migration governance, office hours, a hackathon); coach PMs and engineers.
- Mapped all 275 requests on the intake board into a GenAI taxonomy (summarize / classify / verify / generate): 176 fit none of the four, and the themes underneath—platform enablement, data plumbing, and migration—reframed how leadership prioritized the next scalable builds.
- Product-led a classification and diligence platform for banking benchmarking from discovery to done, authoring the requirements taxonomy engineers built from. Client effort for taxonomy mapping, a primary challenge for 86% of surveyed clients, went from 2.5 weeks → 2 days (analyst effort 5 → 2 days); 1,000 entries 90 min → 12 min in the pilot. Sponsors moved from "build this for us" to co-shaping GenAI strategy.
- Shipped the intake method as code: an agentic skill (transcript → 20+ structured fields, timestamped evidence, confidence scoring, DFV labels, Jira-ready stories, three human approval gates) merged into the firm's skills catalog after peer and security review. Request to prototype: ~4 months → ~1 week.
- Talked the sponsors of six agent requests at the firm's research institute out of building a new platform: one plan on what the firm already ran, sequenced least privilege → the existing AI gateway → a data move (about a dozen of ~3,000 datasets were in use).
- Product lead for governed agentic editorial workflows: five jobs-to-be-done, a PRD per job, binary pass/fail eval rubrics checked by code or a validated LLM judge, release thresholds, a decision log of human rulings, and an AI-interaction training playbook—skills now used on live reports.

- Turned an LLM accuracy plateau into a scoped ML capability: a fine-tuned hierarchical classifier with per-level confidence thresholds and GenAI fallback only on low confidence—deterministic, auditable, and 70–85% fewer LLM calls by design.

Transformation Strategist – McKinsey & Company
Mar 2022 – Apr 2025

Leadership and Operations | Tech Ecosystem | New York, NY

- Redefined mission, operating narrative, and priorities for a 200-person technology function across product, engineering, advisory, and design; spearheaded three organization-level talent-transformation initiatives.
- Brought design thinking, Lean UX, and Agile into Data Labs and the Sustainability practice, shifting teams from output-driven delivery to experimentation and outcomes; embedded Lead Product Strategist for Sustainability roadmap governance.
- Replaced static status updates with interactive leadership narratives that clarified progress, trade-offs, risks, and decisions needed.

Design Director – McKinsey & Company
Nov 2016 – Mar 2022

Group & Portfolio Design Leader | New York, NY

- Redesigned McKinsey's Global IT Support for 45,000+ colleagues from a manual, legacy-heavy service into customer-centric self-service: portal traffic 600/month → 3,000/day, CSAT 4.2 → 4.9, and a smooth shift to remote work in 2020.
- Built a global design-maturity program (assessments, coaching, development paths, intervention playbooks) that moved design from production support to strategic capability across a 4,000-person technology function.
- Managed a global design and research group; hired across New York, Boston, Prague, and New Delhi; led calibration, critique, and leadership development.

SELECTED WORK

Agentic intake skill – transcript to requirements with human gates (2026)
- Nine-step agentic workflow that converts recorded discovery calls into structured intake artifacts and Jira-ready stories with confidence scoring and three approval gates; merged into the firm's skills catalog. Public, from-scratch version with an eval harness (faithfulness, coverage, judge validation): github.com/christine-q-nguyen/transcript-to-requirements.

Classification platform with confidence routing (2025–2026)
- One reusable capability instead of four point tools; semantic + ML + GenAI arbitration with per-level confidence gating; client mapping effort 2.5 weeks → 2 days; three internal teams use it, and it runs as an MCP server.

Governed agentic editorial workflows (2026)
- PRDs, binary eval rubrics, release thresholds, a decision log ("flag, never silently correct"), and a training playbook for a family of drafting and review skills used on live research reports.

Cross-practice GenAI opportunity taxonomy (2026)
- 275 demands classified into summarize / classify / verify / generate with sub-patterns; the missing bucket—platform, data, migration—changed the prioritization conversation.

MISSION-DRIVEN

Operations – Metro Dallas Homeless Alliance
Apr 2008 – Aug 2009

- Helped launch a one-stop-shop homeless services center that reduced surrounding-area crime by 30%; designed service workflows, triaged clients in the field, improved access to housing, healthcare, and employment services.

EDUCATION & DEVELOPMENT

Post Graduate Program, Generative AI for Business Applications
UT Austin (McCombs) with Great Learning · Completed Jun 2026 · GPA 4.33
ML/DL foundations, transformers, prompt engineering, RAG, fine-tuning, agentic workflows, responsible AI, LLMOps, MCP.

Certifications & Teaching
Claude Certified Associate, Architect, and Developer (Foundations, in progress) · Nielsen Norman Group UX Certificate · Certified Agile Leadership I · Cooper Design Leadership · Cooper Service Design · ServiceNow Fundamentals · Columbia University Design Thinking Instructor · SVA Thesis Advisor

M.P.S., Interactive Telecommunications
New York University
Interaction design, physical computing, programming

B.A. with Honors, Arts & Technology
University of Texas at Dallas

Tools
Cursor, Claude Code, MCP (clients and servers built), Claude / OpenAI / Gemini via enterprise AI gateway, Playwright, Git/GitHub, Jira, Confluence, Figma, Miro